

The background is a dark blue gradient with a subtle, abstract pattern of concentric circles and a spiral shape, creating a sense of depth and movement.

# *Prometheus*

## *a practical workshop* v1.1

Tomer Gabel, November 2024

# The Pillars of Observability

1. Logs

2. Metrics

3. Traces

# The Pillars of Observability

~~1. Logs~~

2. Metrics

~~3. Traces~~



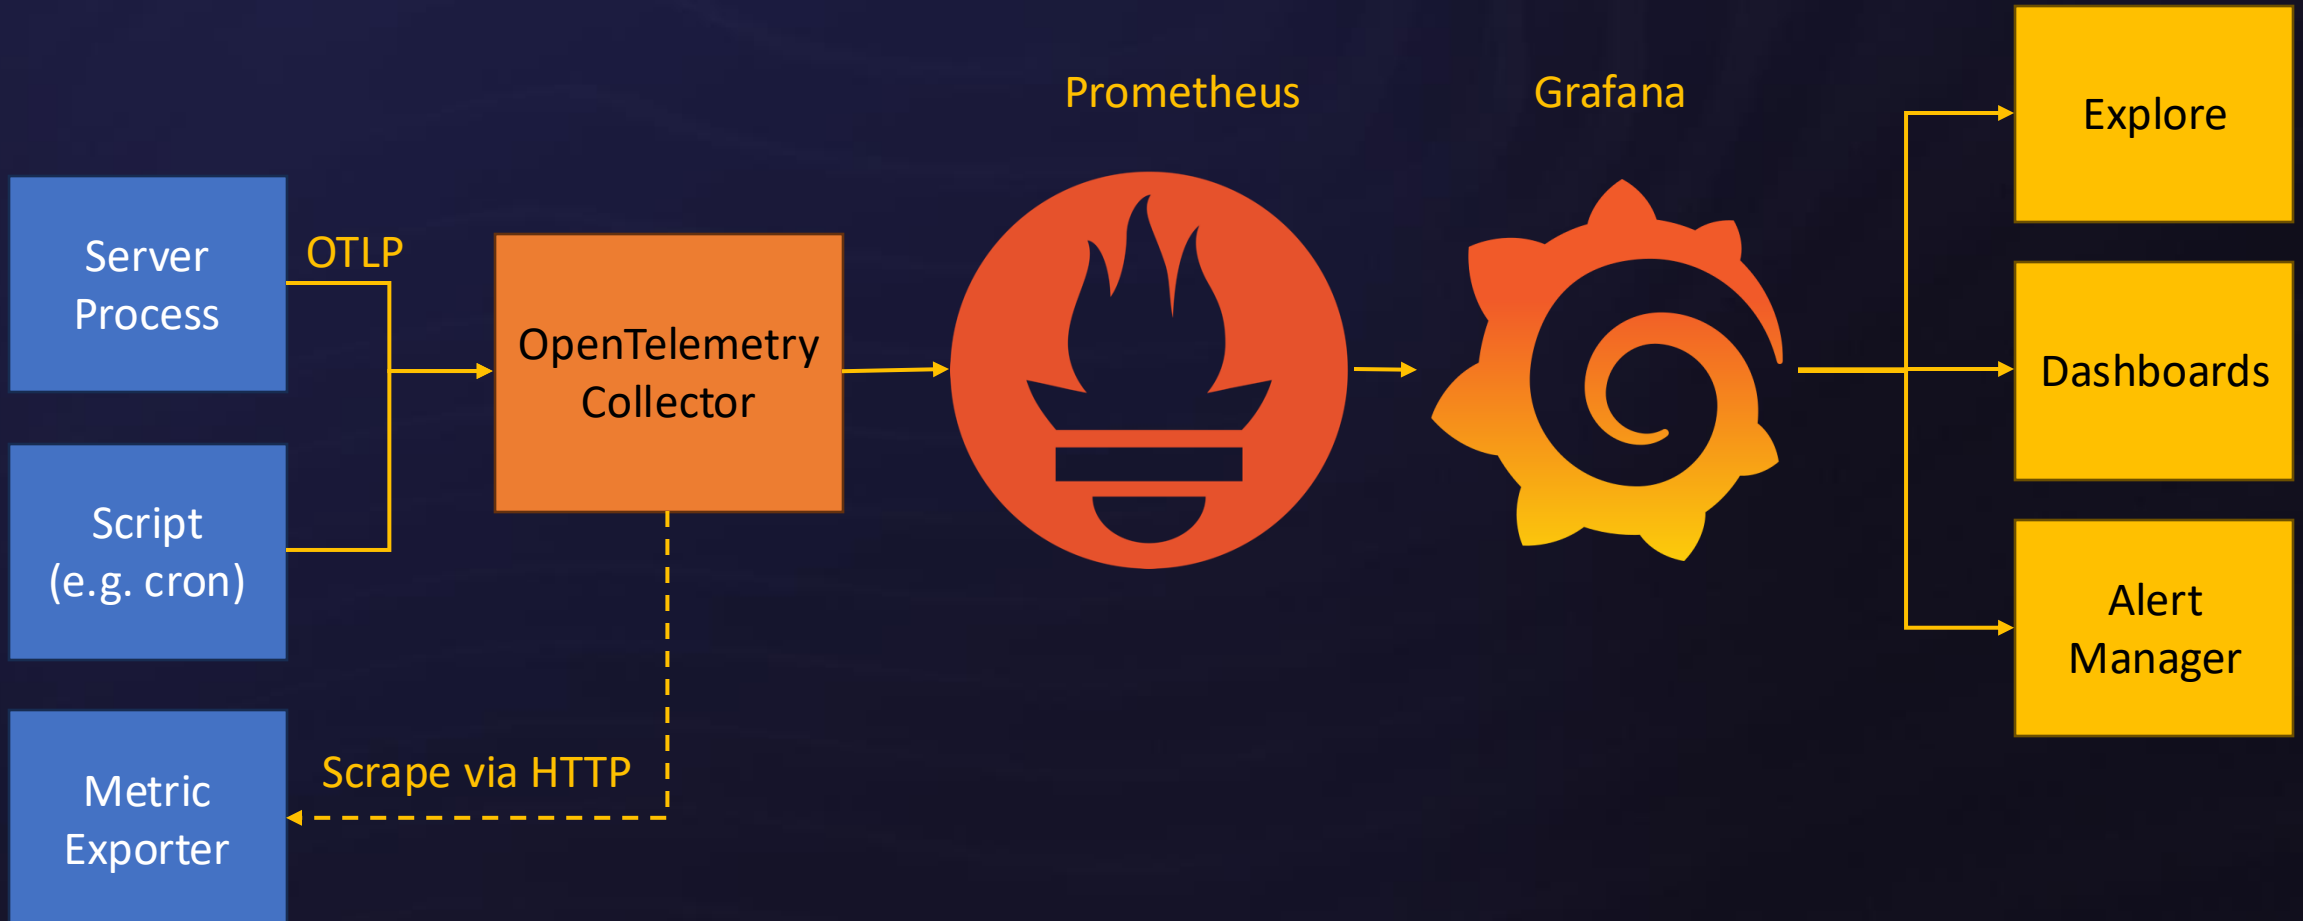
Our focus today.

## Metrics:

Answer *quantitative* questions

1. *How many* requests did I get?
2. *How long* did they take to process?
3. *How much* memory am I using?
4. *How much* free space is left on the disk?

# Enter: Prometheus




# Metrics 101

1. Name

2. Labels

3. Time

4. Value



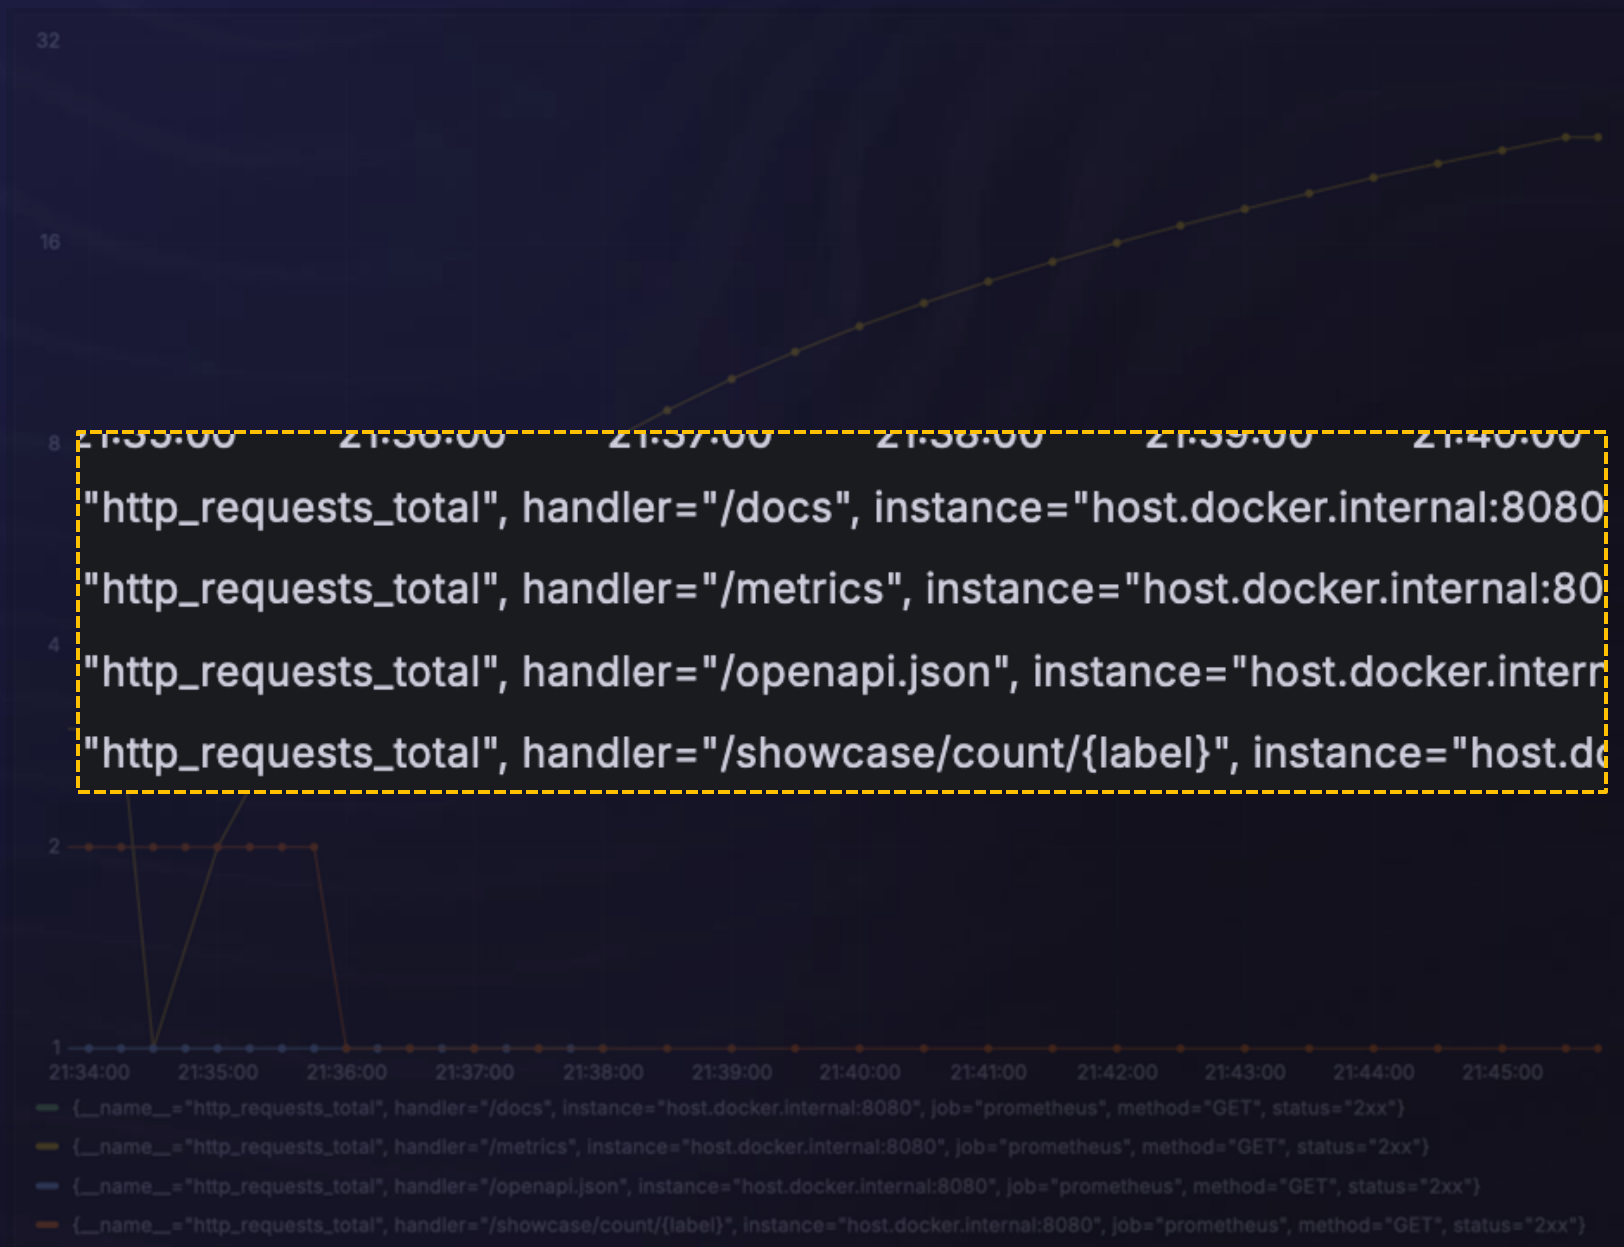
# Metrics 101

1. Name

2. Labels

3. Time

4. Value




# Metrics 101

1. Name

2. Labels

3. Time

4. Value





## A word on cardinality

1. Prometheus is analogous to a *flat database*
  - Partitioned by *time*
  - Labels are like *indexed columns*
2. The DB size is therefore:
  - $\text{Timeframe} * \# \text{ labels} * \# \text{ label values}$

# Prometheus metric types: Counter

1. As the name implies, a metric that *counts*
2. Can go up or down by any amount
3. *Stateful* – value managed by *Prometheus*
4. Examples:
  - # of server requests (by method, path, ...)
  - # of events (“times a user logged in”)

# Counting the Python way


```
counter = meter.create_counter(  
    "my_count",  
    description="Event count"  
)
```

```
@router.get("/count/{label}")
```

```
def count(label: str) -> dict:
```

```
    counter.add(1, {"my_label": label})
```

```
    return {"status": "ok"}
```



# Prometheus metric types: Gauge

1. A metric that is sampled *on demand*
2. Each sample is *independent*
3. Common examples:
  - CPU core temperature in °C
  - Free space on /dev/sda2

# Gauging the Python way

```
gauge = meter.create_gauge(  
    "my_value",  
    description="Some stateful value"  
)
```

```
class GaugeData(BaseModel):
```


```
    value: int
```

```
@router.put("/gauge")
```

```
def gauge_value(data: GaugeData) -> dict:
```

```
    gauge.set(data.value)
```

```
    return {"status": "ok"}
```




# Prometheus metric types: Histograms

1. Some metrics *cannot be represented* with one value
  - “*How long* do requests to /login take?”
  - Do you mean average? Mean? 90th percentile?
2. These values produce a *distribution*
3. Managed with buckets. Latency, for example:
  - 0–10ms, 10–100ms, 100–1000ms, ...

# Histograms the Python way

```
histogram = meter.create_histogram(  
    "my_duration",  
    unit="ms",  
    description="Event duration"  
)
```

```
@router.get("/duration/{length_ms}")  
def duration(length_ms: int) -> dict:  
    histogram.record(length_ms)  
    return {"status": "ok"}
```



# Not all histograms are made equal

*Latencies* vary wildly (ms vs second), as do *sizes* (payload in KB, disk in GB).

*Views* to the rescue!

```
duration_view = View(  
    instrument_name="my_duration",  
    aggregation=ExplicitBucketHistogramAggregation(range(0, 10000, 500)),  
)  
meter_provider = MeterProvider(  
    metric_readers=[oltp_exporter],  
    views=[duration_view],  
)  
set_meter_provider(meter_provider)
```



The background is a dark blue gradient. It features a series of concentric circles on the right side, which transition into a spiral pattern on the left side. The spiral is composed of many thin, overlapping lines that create a sense of depth and movement.

Lab Time

# Showcase

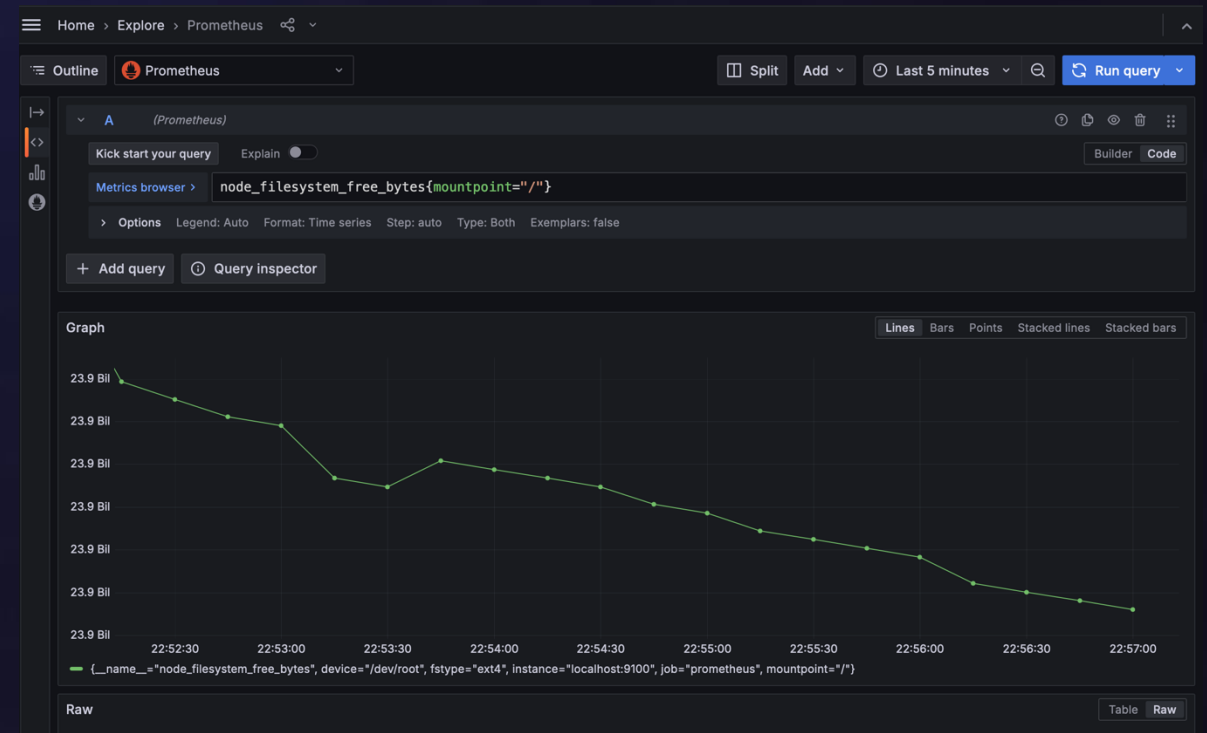
1. *Familiarize yourself* with the lab setup
2. *Get handsy* with Grafana
3. Open lab-showcase from the *class materials*
4. You have *30 minutes* to explore!

The background is a dark blue gradient. It features a large, faint, light blue spiral that starts from the left and curves towards the right. Overlaid on this are several concentric, semi-transparent light blue circles that create a sense of depth and motion.

# Part 2: PromQL

# PromQL: Fear, Terror and Ruthless Efficiency

1. Extensive, powerful
2. ... not entirely trivial



# PromQL: Fear, Terror and Ruthless Efficiency

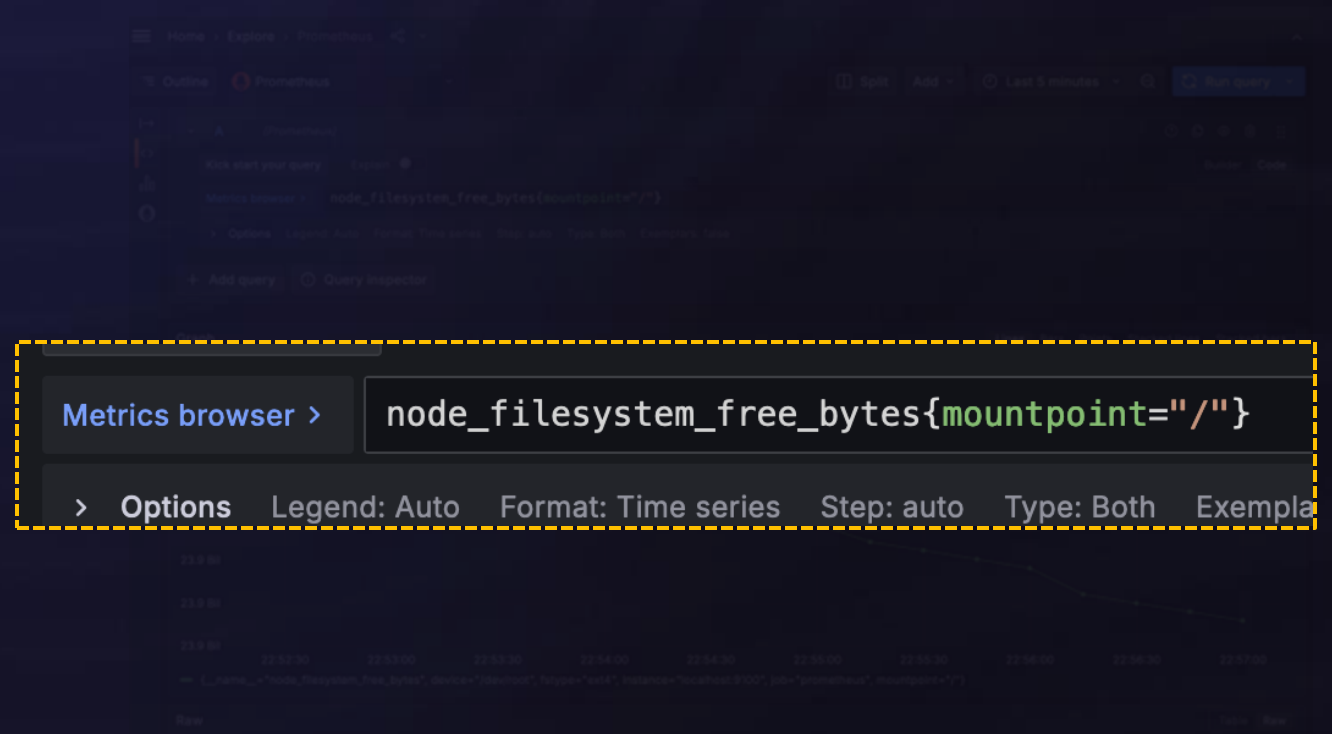
1. Extensive, powerful

2. ... not entirely trivial

3. Basically:

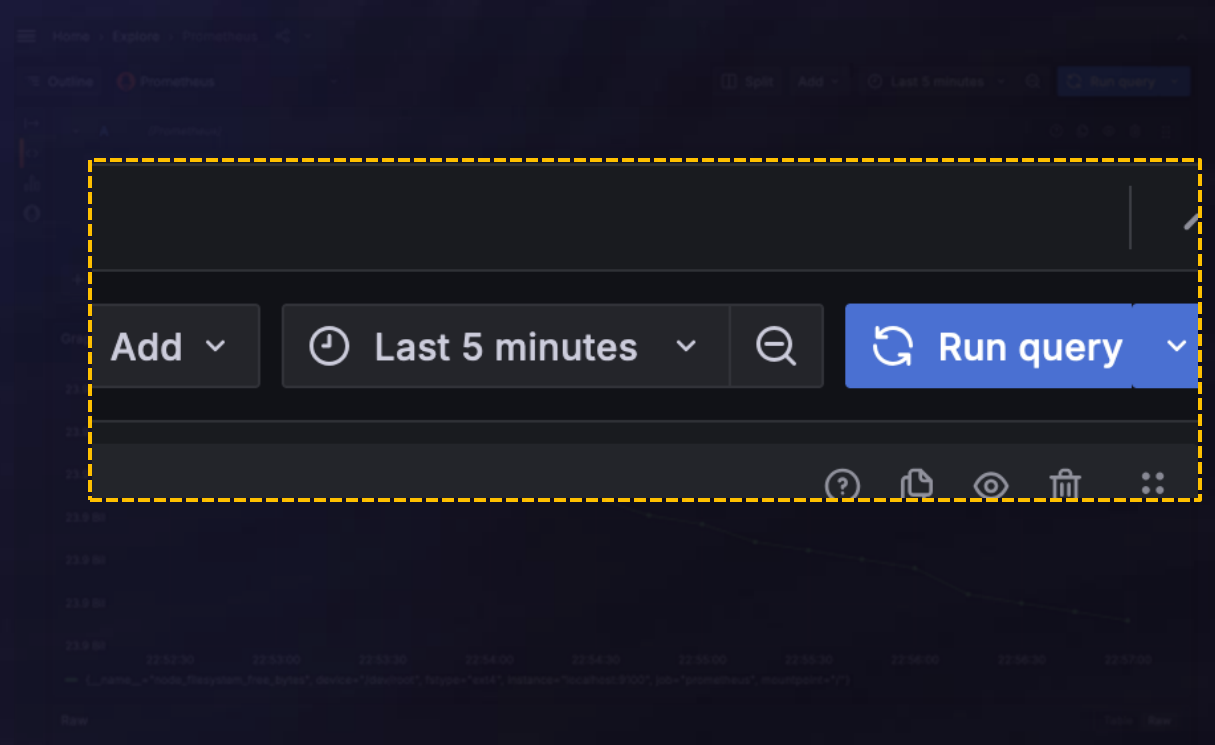
- A *query*

- A *timeframe*



# PromQL: Fear, Terror and Ruthless Efficiency


1. Extensive, powerful
2. ... not entirely trivial
3. Basically:
  - A *query*
  - A *timeframe*



## PromQL: Maths!

1. [Metric, labels, time] yields *one value*
2. Some observations require *multiple* values
3. Expressed with basic arithmetic:
  - % free disk space = *free* disk bytes / *total* disk bytes
  - % CPU = *CPU used seconds* / (# seconds \* *# cores*)

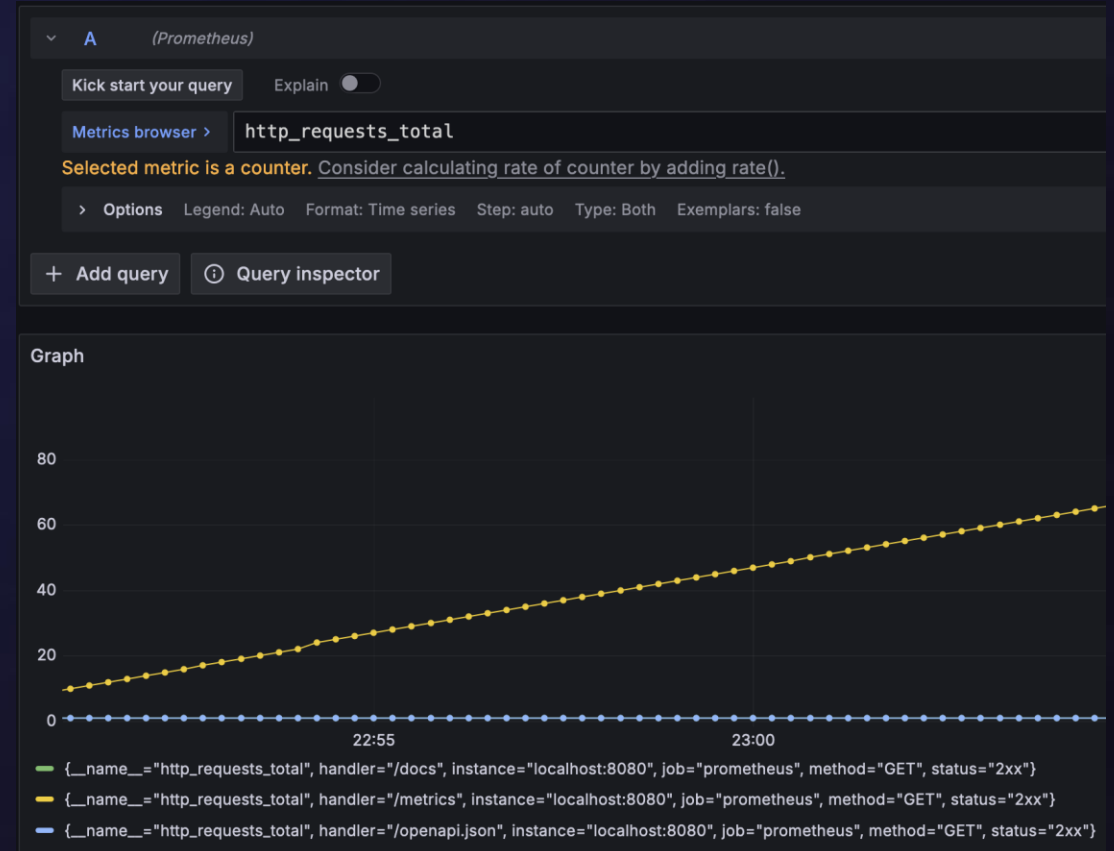
# PromQL: Maths





# PromQL: Rates

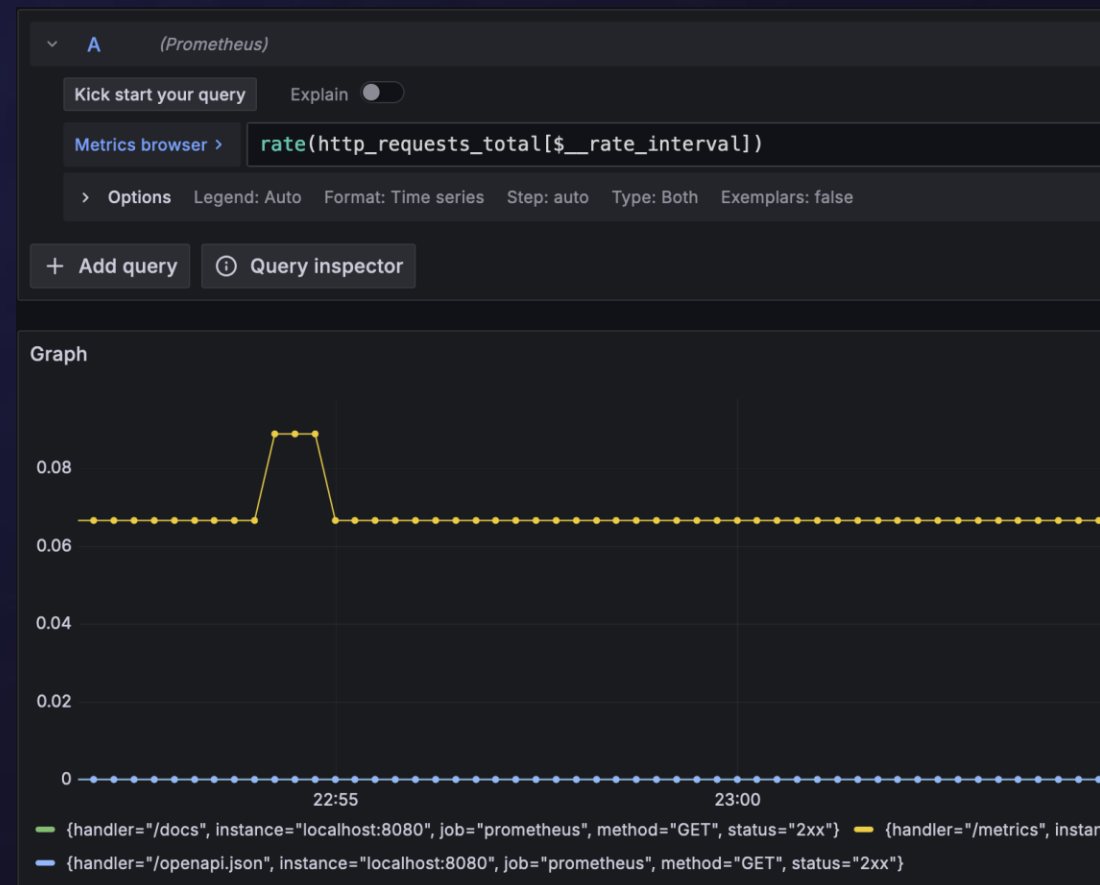
1. Counters are very useful
2. Can't use them *naïvely*



# PromQL: Rates

1. Counters are very useful
2. Can't use them *naïvely*
3. We need a *rate derivation*
4. Magic incantation for the win:

`rate(metric{...}[$__rate_interval])`



# PromQL: Aggregations

1. Consider *http\_server\_duration\_milliseconds\_count*
  - A simple counter
  - Potentially many services, paths
2. Suppose you want the grand total?  
`sum(http_server_duration_milliseconds_count)`

# PromQL: Aggregations

1. Consider *http\_server\_duration\_milliseconds\_count*
  - A simple counter
  - Potentially many services, paths
2. How about the 5 most used paths?  
`topk(5, http_server_duration_milliseconds_count)`

# PromQL: Aggregations

1. What are the top 3 most time-consuming endpoints?

```
topk(3, http_server_duration_milliseconds_sum)
```

2. How many processes run which Python version?

```
count by (major) (python_info)
```

3. What is the highest rate of pagefaults recorded?


```
max(rate(node_vmstat_pgfault[$__rate_interval]))
```

And now,  
what you've all  
been waiting for

# Quantiles (a.k.a. percentiles)


1. A histogram is a set of *counters*
  - 0–100ms, 100–200ms, 200–300ms...
2. What is the 95th *percentile*?
  - Linear regression FTW... but maths!
3. PromQL to the rescue:

*histogram\_quantile(0.99, http\_server\_duration\_milliseconds\_bucket)*




Graph

Lines Bars Points **Stacked lines** Stacked bars



Graph

Lines **Bars** Points Stacked lines Stacked bars





# Key Takeaways

## 1. Instrument early, instrument often

There's no such thing as "too much data"

## 2. Beware high *cardinality*

Aggregate, don't transact (that's what logs/traces are for)

## 3. Learn ye *PromQL*

That's where the real leverage is

The background is a dark blue gradient. It features a series of concentric circles on the right side, which transition into a spiral pattern on the left side. The spiral is composed of many thin, overlapping lines that create a sense of depth and movement.

Lab time, again!

 tomer@substrate.co.il

 @tomerg

 <https://github.com/holograph/prometheus-workshop-service-python>

Thank you for your attention

# Questions?

